

An Overview of Reinforcement Learning

Emmanuel Madrigal
emmadrigal@estudiantec.cr

Abstract

This document provides an overview of reinforcement learning, this is a machine learning technique which deals with the optimization of a reward function when the algorithm can only interact with an intermediary environment. Some general algorithms are shown and the most relevant frameworks are discussed. Its current challenges and some use cases are also presented.




Figure 1. Machine learning system

1. What is it?

Reinforcement learning is learning what to do to maximize a numerical reward signal. The learner is not told which actions to take, but must discover which actions yield the most reward by trying them. In the most interesting and challenging cases, actions may affect not only the immediate reward but also the next situation and, through that, all subsequent rewards. These two characteristics (trial-and-error search and delayed reward) are the two most important distinguishing features of reinforcement learning [1].

Reinforcement learning differs from supervised learning which comprises learning from a training set of labeled examples provided by a knowledgeable external supervisor. Each example is a description of a situation together with a specification—the label—of the correct action the system should take to that situation. The object of supervised learning is for the system to extrapolate, or generalize, its responses so it acts correctly in situations not present in the training set. In interactive problems it is often impractical to get examples of desired behavior that are both correct and representative of all the situations in which the agent has to act. In uncharted territory (where learning to be most beneficial), an agent must be able to learn from its own experience [1].

Reinforcement learning is also different from what machine learning researchers call unsupervised learning, which is typically about finding structure hidden in collections of unlabeled data. Reinforcement learn-

ing is trying to maximize a reward signal instead of trying to find hidden structure. Uncovering structure in an agent's experience can be useful in reinforcement learning, but by itself does not address the reinforcement learning problem of maximizing a reward signal [1].

2. What are its principles?

2.1. Agent and Environment

Reinforcement learning involves an interaction between an active decision-making agent and its environment, within which the agent seeks to achieve a goal despite uncertainty about its environment. The agent's actions may affect the future state of the environment (e.g., the next chess position, the level of reservoirs of the refinery, the robot's next location and the future charge level of its battery), affecting the options and opportunities available to the agent at later times. Correct choice requires taking into account indirect, delayed consequences of actions, and thus may require foresight or planning. The effects of actions cannot be fully predicted; thus the agent must monitor its environment frequently and react appropriately. Use cases for reinforcement learning involve explicit goals in the sense that the agent can judge progress toward its goal based on what it can sense directly [1].

2.2. Policy

A policy defines the learning agent's way of behaving given a situation or state. Roughly, a policy is a mapping from perceived states of the environment to actions to take when in those states. Sometimes the policy may be a simple function or lookup table, whereas in others it may involve extensive computation such as a search process. The policy is the core of a reinforcement learning agent in the sense that it alone suffices to determine behavior. A reward signal defines the goal in a reinforcement learning problem. On each time step, the environment sends to the reinforcement learning agent a single number called the reward. The agent's sole aim is to maximize the total reward it receives over the long run. The reward signal thus defines what are the good and bad events for the agent. They are the immediate and defining features of the problem faced by the agent. The reward signal is the primary basis for altering the policy; if an action selected by the policy results in a low reward, then the policy may change to select some other action in that situation. Reward signals may be stochastic functions of the state of the environment and the actions taken [1].

2.3. Reward and Value

Whereas the reward signal shows what is good in an immediate sense, a value function specifies what is good in the long run. Roughly, the value of a state is the total amount of reward an agent can expect to accumulate over the future, starting from that state. Whereas rewards determine the immediate, intrinsic desirability of environmental states, values show the long-term desirability of states after taking into account the states that are likely to follow, and the rewards available in those states. For example, a state might always yield a low immediate reward but still have a high value because it is regularly followed by other states that yield high rewards. Or the reverse could be true. Rewards are primary, whereas values, as predictions of rewards, are secondary. Without rewards there could be no values, and the only purpose of estimating values is to achieve more reward. It is values with which we are most concerned when making and testing decisions. Action choices are made based on value judgments. We seek actions that bring about states of highest value, not highest reward, because these actions result in the greatest amount of reward for us over the long run. Unfortunately, it is much harder to determine values than it is to determine rewards. Rewards are given directly by the environment, but values must be estimated and re-estimated from the sequences of observations an agent

makes over its entire lifetime. In fact, the most important component of almost all reinforcement learning algorithms we consider is a method for efficiently estimating values. The central role of value estimation is arguably the most important thing we have learned about reinforcement learning over the last few decades [1].

2.4. Environment Model

The fourth and final element of some reinforcement learning systems is a model of the environment. This is something that mimics the behavior of the environment, or that allows inferences to be made about how the environment will behave. For example, given a state and action, the model might predict the resultant next state and next reward. Models are used for planning, by which we mean any way of deciding on a course of action by considering possible future situations before they are actually experienced. Methods for solving reinforcement learning problems that use models and planning are called model-based methods, as opposed to simpler model-free methods that are explicitly trial-and-error learners—viewed as almost the opposite of planning [1].

3. What algorithms exist?

3.1. Monte Carlo

Recall that the value of a state is the expected return—expected cumulative future discounted reward—starting from that state. An obvious way to estimate it from experience, then, is simply to average the returns observed after visits to that state. As more returns are observed, the average should converge to the expected value. This idea underlies all Monte Carlo methods.

In particular, suppose we wish to estimate $v_\pi(s)$, the value of a state under policy π , given a set of episodes obtained by following π and passing through s . Each occurrence of states in an episode is called a visit to s . Of course, s may be visited multiple times in the same episode; let us call the first time it is visited in an episode the first visit t to s . The first-visit MC method estimates $v_\pi(s)$ as the average of the returns following first visits t to s , whereas the every-visit MC method averages the returns following all visits t to s .

As the number of visits increases, the function V converges to $v_\pi(s)$;

Algorithm 1 First-visit MC prediction for estimating $V \approx v_\pi$

```

1: Initialize:
2:  $\pi \leftarrow$  policy to be evaluated
3:  $V \leftarrow$  an arbitrary state-value function
4:  $Returns \leftarrow$  an empty list, for all  $s \in S$ 
5: while True do
6:   Generate an episode using  $\pi$ 
7:   for each state  $s$  in the episode do
8:      $G \leftarrow$  the return that follows the first occurrence of  $s$ 
9:     Append  $G$  to  $Returns(s)$ 
10:     $V(s) \leftarrow average(Returns(s))$ 

```

3.2. Q-Learning

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

The learned action-value function Q directly approximates q_* which is the optimal action-value function [1].

Algorithm 2 Q-learning for estimating $\pi \approx \pi_*$

```

1: Initialize  $Q(s, a)$ , for all  $s \in S, a \in A(s)$ , randomly and  $Q(\text{terminalstate}, \cdot) = 0$ 
2: if Remaining Episodes then
3:   Initialize  $S$ 
4:   if Remaining Steps then
5:     Choose  $A$  from  $S$  using policy derived from  $Q$ 
6:     Take action  $A$ , observe  $R, S'$ 
7:      $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$ 
8:      $S \leftarrow S'$ 

```

3.3. Deep Q-Learning [2]

This is an extension of Q-Learning, that combines it with artificial neural networks. Minh et al showed in [2] how a single reinforcement learning agent can achieve high levels of performance in many different problems without relying on different problem-specific feature sets.

Mnih et al. modified the basic Q-learning procedure in three ways. First, they used a method called experience replay. This method stores the agent's experience at each timestep in a replay memory that is accessed to perform the weight updates. After the game emulator executed action A_t in a state represented by

the image stack S_t , and returned reward R_{t+1} and image stack S_{t+1} , it added the tuple $(S_t, A_t, R_{t+1}, S_{t+1})$ to the replay memory. This memory accumulated experiences over many plays of the same game. At each time step multiple Q-learning updates—a mini-batch—were performed based on experiences sampled uniformly at random from the replay memory. Instead of S_{t+1} becoming the new S_t for the next update as it would in the usual form of Q-learning, a new unconnected experience was drawn from the replay memory to supply data for the next update. Since Q-learning is an off-policy algorithm, it does not need to be applied along connected trajectories.

Mnih et al. used a technique that brought Q-learning closer to the simpler supervised-learning case while still allowing it to bootstrap. Whenever a certain number, C , of updates had been done to the weights w of the action-value network, they inserted the network's current weights into another network and held these duplicate weights fixed for the next C updates of w . The outputs of this duplicate network over the next C updates of w were used as the Q-learning targets. Letting \tilde{q} denote the output of this duplicate network, then instead of (16.3) the update rule was:

$$w_{t+1} = w_t + \alpha [R_{t+1} + \gamma \max_a \tilde{q}(S_{t+1}, a, w_t) - \hat{q}(S_t, A_t, w_t)] \nabla_{w_t} \hat{q}(S_t, A_t, w_t).$$

A final modification of standard Q-learning was also found to improve stability. They clipped the error term $R_{t+1} + \gamma \max_a \tilde{q}(S_{t+1}, a, w_t) - \hat{q}(S_t, A_t, w_t)$ so that it remained in the interval $[-1, 1]$.

3.4. Actor Critic [3]

The principal idea is to split the deep q-learning model in two: one for computing an action based on a state and another one to produce the Q values of the action.

The actor takes as input the state and outputs the best action. It essentially controls how the agent behaves by learning the optimal policy. The critic, on the other hand, evaluates the action by computing the value function. Those two models participate in a game where they both get better in their own role as the time passes. The result is that the overall architecture will learn to play to solve the problem more efficiently than the two methods separately.

The Actor Critic has two main variants: the Asynchronous Advantage Actor Critic (A3C) and the Advantage Actor Critic (A2C).

A3C was introduced in [4]. In essence, A3C implements parallel training where multiple workers in

parallel environments independently update a global value function—hence “asynchronous”. One key benefit of having asynchronous actors is effective and efficient exploration of the state space.

A2C [5] is like A3C but without the asynchronous part; this means a single-worker variant of the A3C. It was empirically found that A2C produces comparable performance to A3C while being more efficient.

3.5. Double DQN [6]

Taking the largest q value (which is noisy) as the best action to take can lead to false positives. If non-optimal actions are regularly given a higher Q value than the optimal best action, the learning will be complicated. The solution is: when we compute the Q target; we use two networks to decouple the action choice from the target Q value generation.

- Use a DQN network to select what is the best action to take for the next state (the action with the highest Q value).
- Use a target network to calculate the target Q value of taking that action at the next state.

3.6. Deep Deterministic Policy Gradient [7]

Deep Deterministic Policy Gradient (DDPG) is an algorithm which concurrently learns a Q -function and a policy. It uses off-policy data and the Bellman equation to learn the Q -function, and uses the Q -function to learn the policy.

This approach is closely connected to Q -learning, and is motivated the same way: if you know the optimal action-value function $Q^*(s, a)$, then in any given state, the optimal action $a^*(s)$ can be found by solving

$$a^*(s) = \arg \max_a Q^*(s, a).$$

DDPG interleaves learning an approximator to $Q^*(s, a)$ with learning an approximator to $a^*(s)$, and it does so in a way which is specifically adapted for environments with continuous action spaces.

Because the action space is continuous, the function $Q^*(s, a)$ is presumed to be differentiable with respect to the action argument. This allows to set up an efficient, gradient-based learning rule for a policy $\mu(s)$ which exploits that fact. Then, instead of running an expensive optimization subroutine each time we wish to compute $\max_a Q(s, a)$, we can approximate it with $\max_a Q(s, a) \approx Q(s, \mu(s))$.

3.7. SARSA [8]

The Sarsa algorithm is an On-Policy algorithm for reinforcement learning. The major difference between it and Q -Learning, is that the maximum reward for the next state is not necessarily used for updating the Q -values. Instead, a new action, and therefore reward, is selected using the same policy that determined the original action. The name Sarsa actually comes from the fact that the updates are done using the quintuple $Q(s, a, r, s', a')$. Where: s, a are the original state and action, r is the reward observed in the following state and s', a' are the new state-action pair.

3.8. Other

Other algorithms include:

- Proximal Policy Optimization [9].
- Trust-Region Policy Optimization [10].
- Deterministic Policy Gradient [11].
- Generative Adversarial Imitation Learning [12].
- Hindsight Experience Relay [13].
- Soft Actor Critic [14].
- Twin Delayed DDPG [15].
- Trust Region Policy Optimization [10].

4. Which library or frameworks are the most relevant today?

4.1. KerasRL

keras-rl[16] implements some state-of-the art deep reinforcement learning algorithms in Python and seamlessly integrates with the deep learning library Keras.

Furthermore, keras-rl works with OpenAI Gym out of the box. This means that evaluating and playing around with different algorithms is easy.

Implements the following algorithms:

- Deep Q Learning
- Double DQN
- Deep Deterministic Policy Gradient
- Continuous DQN
- Cross-Entropy Method
- Dueling Network DQN
- Deep SARSA

4.2. Tensorforce

Tensorforce [17] is an open-source deep reinforcement learning framework, with an emphasis on modularized flexible library design and straightforward usability for applications in research and practice. Tensorforce is built on top of Google's TensorFlow framework and requires Python 3.

- Proximal Policy Optimization
- Trust-Region Policy Optimization
- Deterministic Policy Gradient
- Deep Q-Network
- Double DQN
- Dueling DQN
- Actor-Critic
- Advantage Actor-Critic

4.3. Stable Baselines

Stable Baselines [18], a set of implementations of Reinforcement Learning (RL) algorithms with a common interface, based on OpenAI Baselines. It is focused on simplicity of use and consistency.

- Advantage Actor Critic
- Actor Critic with Experience Replay
- Actor Critic using Kronecker Factored Trust Region
- Deep Deterministic Policy Gradient
- Deep Q Learning
- Generative Adversarial Imitation Learning
- Hindsight Experience Relay
- Proximal Policy Optimization
- Soft Actor Critic
- Twin Delayed DDPG
- Trust Region Policy Optimization

4.4. TF Agents

TF-Agents [19] is a relatively new but up-and-coming tool, again, from Google. TF-Agents, while newer, is typically seen as more robust and mature. It is a framework designed for notebooks and that makes it perfect for trying out various configurations, hyperparameters, or environments.

- Deep Q Learning
- Double DQN
- Twin Delayed DDPG
- REINFORCE
- Proximal Policy Optimization
- Soft Actor Critic

5. What challenges are there?

5.1. Batch Off-line and Off-Policy Training

Many systems cannot be trained on directly and need to learn from fixed logs of the system's behavior. Most times, we are deploying an RL approach to replace a previous control system, and logs from that policy are available. In future training iterations, batches of data will be available from the most recent iteration of the control algorithm. This setup is an off-line and off-policy training regime where the policy needs to be trained from batches of data [20].

For a production system where drops in performance could be very costly, we want to ensure that the new policy improves upon the previous policy. Estimating the policy's performance without running it on the actual system is termed off-policy evaluation. Off-policy evaluation becomes more challenging as the difference between the policies and the resulting state distributions grows [20].

5.2. Learning On the Real System from Limited Samples

Unlike much of the research performed in deep reinforcement learning, actual systems do not have separate training and evaluation environments. All training data comes from the actual system, and the agent cannot have a separate exploration policy during training as its exploratory actions do not come for free. Instead, the agent must perform reasonably well, and act safely throughout learning. For many systems, this means that exploration must be limited, and the resulting data is

low-variance—very little of the state space may be covered in the logs. In addition, since there is often only one instance of the system approaches that instantiate hundreds or thousands of environments to collect more data for distributed training are usually not compatible with this setup [20].

Almost all of these real-world systems are slow-moving, fragile, or expensive enough that the data they produce is costly, and policy learning must be data-efficient. Where there are off-line logs of the system, these might contain nowhere near the amount of data or data coverage that current RL algorithms expect. Learning iterations on an actual system can take a long time, as slower control frequencies might range from 1-hour to multi-month time-steps, and reward horizons could be on the order of months (e.g. online advertisement, drug therapies). Even with higher-frequency control tasks, the learning algorithm needs to learn quickly from potential mistakes without needing to repeat them multiple times before fixing them. Thus, learning on an actual system requires an algorithm to be both sample-efficient and performant in its operation of the system [20] [21].

5.3. Satisfying Safety Constrains

Almost all physical systems can destroy or degrade themselves and their environment. Considering these systems' safety is necessary for controlling them. Safety is important during system operation, but also during exploratory learning phases. These could be safety considerations either of the system itself (limiting system temperatures, contact forces or maintaining minimum battery levels) or of its environment (avoiding dynamic obstacles, limiting end effector velocities). There may exist a fallback watchdog controller, which would takeover if the learned policy violates the safety constraints, but we consider that it should not be explicitly relied upon [20].

5.4. Partial Observability and Non-Stationarity

Almost all real systems where we would want to deploy reinforcement learning are partially observable. For example, on a physical system, we likely do not have observations of the wear and tear on motors or joints, or the amount of buildup in pipes or vents. On systems that interact with users such as recommender systems, we have no observations of the mental state of the users. Often, these partial observabilities appear as non-stationarity (e.g. as a pump's efficiency degrades) or as stochasticity (e.g. as each robot being operated be-

haves differently) [20].

5.5. Unspecified and Multi-Objective Reward Functions

Reinforcement learning frames policy learning through the lens of optimizing a global reward function, yet most systems have multidimensional costs to be minimized. Most times, system or product owners do not have a clear picture of what they want to optimize. When an agent is trained to optimize one metric, other metrics are discovered that also need to be maintained or improved. Thus, a lot of the work on deploying RL to real systems is in formulating the reward function, which may be multidimensional. Because the global reward function is generally a balance of multiple sub-goals (e.g. reducing both time-to-target and energy use), a proper evaluation should explicitly separate the individual components of the reward function to better understand the policy's tradeoffs [20].

5.6. Explainability

Another essential aspect of real systems is that they are owned and operated by humans, who need to be reassured about the controllers' intentions and require insights regarding failure cases. For this reason, policy explainability is important for real-world policies. Especially where the policy might find an alternative and unexpected approach to controlling a system, understanding the longer term intent of the policy is important for obtaining stakeholder buy-in. In the event of policy errors, being able to understand the error's origins a posteriori is essential [20].

5.7. Multi-task Learning

To achieve general AI, an agent should be able to perform many types of tasks, rather than specializing in just one. The core of this challenge is scalability. It should not take 1000 times as many samples or hours of computation time to learn 1000 different tasks than to learn one single task. Instead, an AI agent should build up a library of general knowledge and learn general skills that can be used across a variety of tasks. Besides being scalable, an AI that learns general skills can also quickly adapt to unfamiliar tasks, enabling consistent performance in dynamic environments [21].

5.8. Learning to Remember

For many real-world tasks, an observation only captures a minor part of the full environment state that

determines the best action. In such partially observable environments, an agent has to take into account not just the current observation, but also past observations to determine the best action.

For example, consider an intelligent agent in the workplace that helps a company support team employee carrying out actions to help address a customer issue. The human employee may ask a customer about a billing issue. That customer might have a home phone, a mobile and an internet account. If the human asks “What’s the outstanding balance on the account?” the agent must remember the course of the conversation to understand which account the human refers to.

Remembering everything in a conversation, however, makes learning a good policy intractable. As humans speak we move from topic-to-topic, changing the subject and looping back again. Some information is very important whereas other information is more tangential. Hence, the challenge is to learn a compact representation that only stores the most salient information [21].

6. Use Cases

6.1. Resources Management in Computer Clusters

Designing algorithms to allocate limited resources to different tasks is challenging and requires human-generated heuristics. In [22] the authors showed how to use RL to automatically learn to allocate and schedule computer resources to waiting jobs, with the objective to minimize the average job slowdown.

They formulated the state space as the current resources allocation and the resources profile of jobs. For action space, they used a trick to allow the agent to choose more than one action at each time step. Reward was the sum of $-1/\text{duration_of_the_job}$ over all the jobs in the system. Then they combined REINFORCE algorithm and baseline value to calculate the policy gradients and find the best policy parameters that give the probability distribution of actions to minimize the objective. See this repository for the implementation.

6.2. Traffic Light Control

In [23], researchers tried to design a traffic light controller to solve the congestion problem. Tested only on simulated environment though, their methods showed superior results than traditional methods and shed a light on the potential uses of multi-agent RL in designing traffic system.

They put five agents in the five-intersection traf-

fic network, with an agent at the central intersection to control traffic signaling. The state was defined as an eight-dimensional vector with each element representing the relative traffic flow of each lane. Eight choices were available to the agent, each representing a phase combination, and they defined the reward function as a reduction in the delay compared with earlier time step. The authors used DQN to learn the Q value of the state, action pairs.

6.3. Robotics

Robotics is an area where reinforcement learning has shown a lot of promise. A general survey can be found in [24]. In particular, [25] trained a robot to learn policies to map raw video images to the robot’s actions. They fed the RGB images to a CNN and outputs were the motor torques. The RL component was the guided policy search to generate training data that came from its own state distribution.

6.4. Personalized Recommendations

Previous work of news recommendations faced several challenges including the changing dynamic of news, users get bored and Click Through Rate cannot show the retention rate of users. Guanjie et al. have applied RL in news recommendation system to combat the problems in [26].

They constructed four categories of features:

- User features
- Context features as the state features of the environment
- User-news features
- News features as the action features.

The four features were input to the Deep Q-Network (DQN) to calculate the Q-value. The agent chose a list of news to recommend based on the Q-value, and the user’s click on the news was a part of the reward the RL agent received.

6.5. Bidding and Advertising

Researchers from Alibaba Group published [27] and claimed that their distributed cluster-based multi-agent bidding solution (DCMAB) has achieved promising results and thus they plan to conduct a live test in Taobao platform.

Taobao ad platform is a place for merchants to place a bid to display ad to the customers. This could

be a multi-agent problem because the merchants are bidding against each other and their actions are inter-related. In the paper, merchants and customers were clustered into different groups to reduce computational complexity. The state space of the agents showed the cost-revenue status of the agents, action space was the bid (continuous), and reward was the revenue caused by the customer cluster.

6.6. Games

RL is so well known these days because it is the mainstream algorithm used to solve different games and sometimes achieve superhuman performance.

The most famous one must be AlphaGo [28] and AlphaGo Zero [29]. AlphaGo, trained with countless human games, already achieved superhuman performance by using value network and Monte Carlo tree search (MCTS) in its policy network. Yet, the researchers later on thought back and tried a purer RL approach — train it from scratch. The researchers let the new agent, AlphaGo Zero, played with itself and beat AlphaGo 100–0.

6.7. Deep Learning

More and more attempts to combine RL and other deep learning architecture can be seen, and they showed impressive results.

One of the most influential work in RL is the pioneering work of DeepMind to combine CNN with RL [2]. By doing so, the agent can “see” the environment through high-dimensional sensory and then learn to interact with it.

RL and RNN is another combinations people used to try a novel idea. RNN is a neural network that has “memories”. When joint with RL, RNN gives the agents’ ability to memorize things. For example, [30] joint LSTM with RL to create Deep Recurrent Q-Network (DRQN) to play Atari 2600 games. In [31] the authors also used RNN and RL to solve chemical reaction optimization problem.

DeepMind showed [32] how to use generative models and RL to generate programs. In the model, the adversarially trained agent used the signal as rewards to improve the actions, instead of propagating the gradients to the input space as in the GAN training.

References

[1] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*. The MIT Press, 2 ed., 2017.

[2] A. Nair, P. Srinivasan, S. Blackwell, C. Alceick, R. Fearon, A. D. Maria, V. Panneerselvam, M. Suleyman, C. Beattie, S. Petersen, S. Legg, V. Mnih, K. Kavukcuoglu, and D. Silver, “Massively parallel methods for deep reinforcement learning,” *CoRR*, vol. abs/1507.04296, 2015.

[3] V. R. Konda and J. N. Tsitsiklis, “Actor-critic algorithms,” in *Advances in neural information processing systems*, pp. 1008–1014, 2000.

[4] V. Mnih, A. P. Badia, M. Mirza, A. Graves, T. Lillicrap, T. Harley, D. Silver, and K. Kavukcuoglu, “Asynchronous methods for deep reinforcement learning,” in *International conference on machine learning*, pp. 1928–1937, 2016.

[5] Y. Wu, E. Mansimov, S. Liao, A. Radford, and J. Schulman, “Openai baselines: Acktr & a2c,” Aug 2017.

[6] H. V. Hasselt, “Double q-learning,” in *Advances in Neural Information Processing Systems 23* (J. D. Lafferty, C. K. I. Williams, J. Shawe-Taylor, R. S. Zemel, and A. Culotta, eds.), pp. 2613–2621, Curran Associates, Inc., 2010.

[7] T. P. Lillicrap, J. J. Hunt, A. Pritzel, N. Heess, T. Erez, Y. Tassa, D. Silver, and D. Wierstra, “Continuous control with deep reinforcement learning,” *arXiv preprint arXiv:1509.02971*, 2015.

[8] G. A. Rummery and M. Niranjan, *On-line Q-learning using connectionist systems*, vol. 37. University of Cambridge, Department of Engineering Cambridge, UK, 1994.

[9] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” *arXiv preprint arXiv:1707.06347*, 2017.

[10] J. Schulman, S. Levine, P. Abbeel, M. Jordan, and P. Moritz, “Trust region policy optimization,” in *International conference on machine learning*, pp. 1889–1897, 2015.

[11] D. Silver, G. Lever, N. Heess, T. Degris, D. Wierstra, and M. Riedmiller, “Deterministic policy gradient algorithms,” 2014.

[12] J. Ho and S. Ermon, “Generative adversarial imitation learning,” in *Advances in neural information processing systems*, pp. 4565–4573, 2016.

[13] M. Andrychowicz, F. Wolski, A. Ray, J. Schneider, R. Fong, P. Welinder, B. McGrew, J. Tobin, P. Abbeel, and W. Zaremba, “Hindsight experience replay,” *CoRR*, vol. abs/1707.01495, 2017.

[14] T. Haarnoja, A. Zhou, P. Abbeel, and S. Levine, “Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor,” *CoRR*, vol. abs/1801.01290, 2018.

[15] S. Dankwa and W. Zheng, “Twin-delayed ddpg: A deep reinforcement learning technique to model a continuous movement of an intelligent robot agent,” in *Proceedings of the 3rd International Conference on Vision, Image and Signal Processing*, pp. 1–5, 2019.

[16] M. Plappert, “keras-rl,” <https://github.com/keras-rl/keras-rl>, 2016.

[17] A. Kuhnle, M. Schaarschmidt, and K. Fricke, “Tensor-

- force: a tensorflow library for applied reinforcement learning." Web page, 2017.
- [18] A. Hill, A. Raffin, M. Ernestus, A. Gleave, A. Kanervisto, R. Traore, P. Dhariwal, C. Hesse, O. Klimov, A. Nichol, M. Plappert, A. Radford, J. Schulman, S. Sidor, and Y. Wu, "Stable baselines." <https://github.com/hill-a/stable-baselines>, 2018.
 - [19] Sergio Guadarrama, Anoop Korattikara, Oscar Ramirez, Pablo Castro, Ethan Holly, Sam Fishman, Ke Wang, Ekaterina Gonina, Neal Wu, Efi Kokiopoulou, Luciano Sbaiz, Jamie Smith, Gábor Bartók, Jesse Berent, Chris Harris, Vincent Vanhoucke, Eugene Brevdo, "TF-Agents: A library for reinforcement learning in tensorflow." <https://github.com/tensorflow/agents>, 2018. [Online; accessed 25-June-2019].
 - [20] A. Gabriel Dulac, D. J. Mankowitz, and T. Hester, "Challenges of real-world reinforcement learning," *CoRR*, 2019.
 - [21] "The next challenges for reinforcement learning," Jan 2018.
 - [22] H. Mao, M. Alizadeh, I. Menache, and S. Kandula, "Resource management with deep reinforcement learning," in *Proceedings of the 15th ACM Workshop on Hot Topics in Networks*, pp. 50–56, 2016.
 - [23] I. Arel, C. Liu, T. Urbanik, and A. G. Kohls, "Reinforcement learning-based multi-agent system for network traffic signal control," *IET Intelligent Transport Systems*, vol. 4, no. 2, pp. 128–135, 2010.
 - [24] J. Kober, J. A. Bagnell, and J. Peters, "Reinforcement learning in robotics: A survey," *The International Journal of Robotics Research*, vol. 32, no. 11, pp. 1238–1274, 2013.
 - [25] S. Levine, C. Finn, T. Darrell, and P. Abbeel, "End-to-end training of deep visuomotor policies," *The Journal of Machine Learning Research*, vol. 17, no. 1, pp. 1334–1373, 2016.
 - [26] G. Zheng, F. Zhang, Z. Zheng, Y. Xiang, N. J. Yuan, X. Xie, and Z. Li, "Drn: A deep reinforcement learning framework for news recommendation," in *Proceedings of the 2018 World Wide Web Conference*, pp. 167–176, 2018.
 - [27] J. Jin, C. Song, H. Li, K. Gai, J. Wang, and W. Zhang, "Real-time bidding with multi-agent reinforcement learning in display advertising," in *Proceedings of the 27th ACM International Conference on Information and Knowledge Management*, pp. 2193–2201, 2018.
 - [28] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. Van Den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, *et al.*, "Mastering the game of go with deep neural networks and tree search," *nature*, vol. 529, no. 7587, pp. 484–489, 2016.
 - [29] D. Silver and D. Hassabis, "Alphago zero: Starting from scratch," *London, United Kingdom: Deepmind*, 2017.
 - [30] M. J. Hausknecht and P. Stone, "Deep recurrent q-learning for partially observable mdps," *CoRR*, vol. abs/1507.06527, 2015.
 - [31] Z. Zhou, X. Li, and R. N. Zare, "Optimizing chemical reactions with deep reinforcement learning," *ACS central science*, vol. 3, no. 12, pp. 1337–1344, 2017.
 - [32] Y. Ganin, T. Kulkarni, I. Babuschkin, S. M. A. Eslami, and O. Vinyals, "Synthesizing programs for images using reinforced adversarial learning," *CoRR*, vol. abs/1804.01118, 2018.